

The Camera Sensor

A. Giavaras

Contents

1	Camera Sensor	2
1.1	The Pinhole Camera Model	2
1.2	Summary	4
1.3	Questions	4
1.4	Assignments	4
2	Camera Projective Geometry	4
2.1	Problem Definition	4
2.2	Mathematical Formulation	5
2.3	Summary	7
2.4	Questions	7
2.5	Assignments	7
3	Camera Calibration	7
3.1	The Camera Calibration Problem	8
3.2	RQ Factorization	9
3.3	Summary	10
3.4	Questions	10
3.5	Assignments	10
4	Visual Depth Perception	10
4.1	Summary	13
4.2	Questions	13
4.3	Assignments	13
5	Compute Disparity	13
5.1	Disparity Computation Algorithm	14
5.2	Summary	16
5.3	Questions	16
5.4	Assignments	16

1 Camera Sensor

The camera sensor is one of the primary sensors in a vehicle's sensor suite. This is because the camera is a rich sensor that captures incredible detail about the environment around the vehicle. However, it requires extensive processing to make use of the information that is available in that image.

In this section, we will highlight why the camera is a critical sensor for autonomous driving. We will then briefly introduce the concept of image formation and present the pinhole camera model which captures the essential elements of how a camera works in a simple and elegant manner.

Of all the common self-driving car sensors, the camera is the sensor that provides the most detailed appearance information from objects in the environment. Appearance information is particularly useful for scene understanding tasks such as object detection, segmentation and identification. Appearance information is what allows us to distinguish between road signs or traffic lights states, to track turn signals and resolve overlapping vehicles into separate instances. Because of its high resolution output, the camera is able to collect and provide orders of magnitude, more information than other sensors used in self-driving while still being relatively inexpensive. The combination of high valued appearance information and low cost make the camera an essential component of our sensor suite.

1.1 The Pinhole Camera Model

Let us see how the camera manages to collect this huge amount of information. A camera is a passive external receptive sensor. It uses an imaging sensor to capture information conveyed by light rays emitted from objects in the world. This was originally done with film but nowadays we use rather sophisticated silicon chips to gather this information. Light is reflected from every point on an object in all directions, and a portion of these rays travel towards the camera sensor.

Look at the car's reflected rays collected by our imaging surface. Do you think we will get a good representation of the car on the image sensor from this ray-pattern? Unfortunately, no. Using this basic open sensor camera design, we will end up with blurry images because our imaging sensor is collecting light rays from multiple points on the object at the same location on the sensor. The solution to our problem is to put a barrier in front of the imaging sensor with a tiny hole or aperture in its center. The barrier allows only a small number of light rays to pass through the aperture, reducing the blurriness of the image. This model is called the **pinhole camera model** and describes the relationship between a point in the world and its corresponding projection on the image plane, see for instance [1].

The two most important parameters in a pinhole camera model are the distance between the pinhole and the image plane which we call the focal length and is typically denoted with f . The focal length defines the size of the object projected onto the image and plays an important role in the camera focus when using lenses to improve camera performance.

Remark 1.1. Focal Length f

Specifically, we define the focal length f as the distance between the camera and the image coordinate frames along the z -axis of the camera coordinate frame.

The coordinates of the center of the pinhole, (c_u, c_v) , which we call the camera center, these coordinates to find the location on the imaging sensor that the object projection will inhabit.

Although the pinhole camera model is very simple, it works surprisingly well for representing the image creation process. By identifying the focal length and the camera's center for a specific camera configuration, we can mathematically describe the location that a ray of light emanating from an object in the world will strike the image plane. This allows us to form a measurement model of image formation for use in state estimation and object detection.

Remark 1.2. Some History

A historical example of the pinhole camera model is the camera obscura, which translates to dark room camera in English. Historical evidence shows that this form of imaging was discovered as early as 470 BC in ancient China and Greece. It's simple construction with a pinhole aperture in front of an imaging surface makes it easy to recreate on your own, and is in fact a safe way to watch solar eclipse if you're so inclined.

Nowadays cameras allow us to collect extremely high resolution data. They can operate in low-light conditions or at a long range due to the advanced lens optics that gather a large amount of light and focus it accurately on the image plane. The resolution and sensitivity of camera sensors continues to improve, making cameras one of the most ubiquitous sensors on the planet; think for example how many cameras you may be owning.

These advances are also extremely beneficial for understanding the environment around a self-driving car. Cameras specifically designed for autonomous vehicles need to work well in a wide range of lighting conditions and in distances to objects. These properties are essential to driving safely in all operating conditions.

1.2 Summary

This was an introductory section. We discussed the usefulness of the camera as a sensor for autonomous driving. We also saw the pinhole camera model in its most basic form, which we'll use to construct algorithms for visual perception. In the next section, we will describe how an image is formed, a process referred to as projective geometry, which relates objects in the world to their projections on the imaging sensor.

1.3 Questions

1. Describe the pinhole camera model. Why is it useful?

1.4 Assignments

1. Using OpenCV read and display an image.

2 Camera Projective Geometry

In this section, you will learn how to model the camera's projective geometry through the coordinate system transformation. These transformations can be used to project points from the world frame to the image frame, building on the pinhole camera model from section 1.1.

You will then model these transformations using matrix algebra and apply them to a 3D point to get its 2D projection onto the image plane. Finally, you will learn how camera 2D images are represented in software. Equipped with the projection equations in image definitions, you will then be able to create algorithms for detecting objects in 3D and localizing the self-driving car later on in the course.

2.1 Problem Definition

First, let's define the problem we need to solve. Let's start with a point \mathbf{O}_{world} defined at a particular location in the world coordinate frame. We want to project this point from the world frame to the camera image plane. Light travels from the \mathbf{O}_{world} on the object through the camera aperture to the sensor surface. You can see that our projection onto the sensor surface through the aperture results in flipped images of the objects in the world. To avoid this confusion, we usually define a virtual image plane in front of the camera center. Let's redraw our camera model with this sensor plane instead of the real image plane behind the camera lens. We will call this model the simplified camera model, and need

to develop a model for how to project a point from the world frame coordinates X_w, Y_w, Z_w to, image coordinates u, v .

We begin by defining the following characteristics of the cameras that are relevant to our problem. First, we select a world frame, W , in which to define the coordinates of all objects and the camera. We also define the camera coordinate frame, C , as the coordinate frame attached to the center of our lens aperture known as the **optical sensor**. We can define a translation vector and a rotation matrix to model any transformation between a world coordinate frame W and another, and in this case, we'll use the world coordinate frame and the camera coordinate frame. We refer to the parameters of the camera pose as the extrinsic parameters, as they are external to the camera and specific to the location of the camera in the world coordinate frame. We define our image coordinate frame as the coordinate frame attached to our virtual image plane emanating from the optical center. The image pixel coordinate system however, is attached to the top left corner of the virtual image plane. So we need to adjust the pixel locations to the image coordinate frame. Next, we define the focal length f as the distance between the camera and the image coordinate frames along the z_C -axis of the camera coordinate frame. Finally, our projection problem reduces to two steps.

1. Project from the world to the camera coordinates
2. Project from the camera coordinates to the image coordinates

We can then transform image coordinates to pixel coordinates through scaling and offset. We now have the geometric model to allow us to project a point from that world frame to the image coordinate frame, whenever we want.

2.2 Mathematical Formulation

Let us formulate the mathematical tools needed to perform this projection using linear algebra. First, we begin with the transformation from the world to the camera coordinate frame, namely $W \rightarrow C$. This is performed using the rigid body transformation matrix \mathbf{T} , which has \mathbf{R} and t in it. Hence, we have

$$\mathbf{O}_{camera} = [\mathbf{R}|t]\mathbf{O}_{world} \quad (1)$$

The next step is to transform camera coordinates to image coordinates. To perform this transformation, we define the matrix \mathbf{K} as the three-by-three matrix given by equation 2.

$$\mathbf{K} = \begin{bmatrix} f & 0 & u_0 \\ 0 & f & v_0 \\ 0 & 0 & 1 \end{bmatrix} \quad (2)$$

This matrix depends on camera intrinsic parameters, which means it depends on components internal to the camera such as the camera geometry and the camera lens characteristics. Since both transformations are just matrix multiplications, we can define a matrix \mathbf{P}

$$\mathbf{P} = \mathbf{K}[\mathbf{R}|t] \quad (3)$$

This matrix transforms from the world coordinate frame all the way to the image coordinate frame. The coordinates of point \mathbf{O}_{world} can now be projected to the image plane via the equation

$$\mathbf{O}_{image} = \mathbf{P}\mathbf{O}_{world} \quad (4)$$

So, let's see what we're still missing to compute this equation. When we expect the matrix dimensions, we noticed that the matrix multiplication cannot be performed. In order to remedy this problem, we transform the coordinates of the point \mathbf{O} into homogeneous coordinates. This is done by adding a one at the end of the 3D coordinates.

Remark 2.1. Homogeneous Coordinates

The point geometric primitive can be represented using homogeneous coordinates $\tilde{\mathbf{x}}$. Consider for example a 2D point $\mathbf{x} = (x_1, x_2)$ this can be written as $\tilde{\mathbf{x}} = (\tilde{x}_1, \tilde{x}_2, \tilde{w}) \in P^2$, where vectors that differ only by scale are considered to be equivalent. The $P^2 = R^3 - (0, 0, 0)$ is called the 2D projective space. A homogeneous vector can be converted back into an inhomogeneous vector by dividing through the last element \tilde{w} :

$$\tilde{\mathbf{x}} = (\tilde{x}_1, \tilde{x}_2, \tilde{w}) = \tilde{w}(x_1, x_2, 1) = \tilde{w}\mathbf{x} \quad (5)$$

Homogeneous [points whose last element is $\tilde{w} = 0$ are called ideal points or points at infinity and do not have an equivalent inhomogeneous representation.

So, now the dimensions work and we are all ready to start computing our projections. Now, we need to perform the final step, transforming the image coordinates to pixel coordinates. We do so by dividing x and y by z to get homogeneous coordinates in the image plane.

This is the basic camera projection model. In practice, we usually model more complex phenomena such as non-square pixels, camera access skew, distortion and non unit aspect ratio. Luckily, this only changes the camera \mathbf{K} matrix, and the equations above can be used as is with a few additional parameters.

Now that we have formulated the coordinates of projection of a 3D point onto the 2D image plane, we want to define what values go into the coordinates in a

2D color image. We will start with a grayscale image. We first define a width N and a height M of an image, as the number of rows and columns the image has. Each point in 3D projects to a pixel on the image defined by the u, v coordinates we derived earlier. Zooming in, we can see these pixels is a grid. In grayscale, brightness information is written in each pixel as an unsigned eight bit integer. Some cameras can produce unsigned 16-bit integers for better quality images. For color images, we have a third dimension of value three we call depth. Each channel of this depth represents how much of a certain color exists in the image.

Many other color representations are available, but we will be using the RGB representation, so red green and blue.

In conclusion, an image is represented digitally as an $M \times N \times 3$ array of pixels, with each pixel representing the projection of a 3D point onto the 2D image plane.

2.3 Summary

So, in this section, we discussed how to project 3D points in the world coordinate frame to 2D points in the image coordinate frame. You saw that the equations that perform this projection rely on camera intrinsic parameters as well as on the location of the camera in the world coordinate frame.

This projection model is used in every visual perception algorithm we develop, from object detection to derivable space estimation. Finally, we saw that images are represented in software as an array representing pixel locations. In the next section, we will discuss how to tailor the camera model to a specific camera by computing its intrinsic and extrinsic camera parameters through a process known as camera calibration.

2.4 Questions

2.5 Assignments

3 Camera Calibration

Section 2.1 discussed which camera parameters are needed for projective geometry to work. In this section, we will learn how to get these camera parameters using the mathematical tools of calibration, see [2].

Thus, camera calibration is the process of estimating intrinsic and/or extrinsic parameters. Intrinsic parameters deal with the camera's internal characteristics, such as, its focal length, skew, distortion, and image center. Extrinsic parameters describe its position and orientation in the world. Knowing intrinsic parameters is an essential first step for 3D computer vision, as it allows you

to estimate the scene's structure in Euclidean space and removes lens distortion, which degrades accuracy, [2].

Recall the projection equations. The homogeneous coordinates of point \mathbf{O} in 3D space can be transformed to the camera plane, with the camera projection matrix \mathbf{P} , which includes both extrinsic and extrinsic parameters.

$$\mathbf{O}_{image} = \mathbf{P}\mathbf{O}_{world} \quad (6)$$

Remember, the projected coordinates need to be converted to a homogeneous form to get the u, v pixel locations in pixel coordinates. We can do this by dividing the image coordinates by the z -component. Finally, u and v can then be multiplied with an arbitrary scale s . We multiply by s , as it will be useful later on when we formulate the calibration problem. It is important to note that scale plays a challenging role in understanding monocular image information, as once points are projected from the 3D world onto the 2D image plane, scale information is lost. Points in 3D space along a ray from the camera center, all project to the same location on the image plane, and it is therefore not possible to directly associate a depth to a point given only image information.

3.1 The Camera Calibration Problem

The camera calibration problem is defined as finding these unknown intrinsic and extrinsic camera parameters, shown here in red given n known 3D point coordinates and their corresponding projection to the image plane.

$$\begin{bmatrix} su \\ sv \\ s \end{bmatrix} = \begin{bmatrix} ? & ? & ? & ? \\ ? & ? & ? & ? \\ ? & ? & ? & ? \end{bmatrix} = \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} \quad (7)$$

Our approach will comprise of getting the \mathbf{P} matrix first and then decomposing it into the intrinsic parameters \mathbf{K} and the extrinsic rotation parameters \mathbf{R} and translation parameters \mathbf{t} . For calibration, we use a scene with known geometry to get the location of our 3D points from the 2D image, resolving the scale issue by measuring the actual 3D distance between the points that are observed in the image. The most commonly used example would be a 3D checkerboard, with squares of known size providing a map of fixed point locations to observe. We define our world coordinate frame, in yellow and compute our 3D point coordinates and their projections in the image. Associating 3D points to 2D projections can be done either manually, by clicking on the purple points, for example or automatically, with checkerboard detectors. We can then set up a system of equations to solve for the unknown parameters of \mathbf{P} .

Now, let us form the system of linear equations that needs to be solved. First, we expand the projection equations to three equations through matrix multiplication.

$$su = p_{11}X + p_{12}Y + p_{13}Z + p_{14} \quad (8)$$

$$sv = p_{21}X + p_{22}Y + p_{23}Z + p_{24} \quad (9)$$

$$s = p_{31}X + p_{32}Y + p_{33}Z + p_{34} \quad (10)$$

To get zero on the right-hand side of these equations, we move the right hand side to the left-hand side for each one. Then, we substitute the third equation into equations one and two, and end up with two equations per point. Therefore, if we have n points, we have $2n$ associated equations. Putting these equations in matrix form gives us the following homogeneous linear system.

Since this is a homogeneous linear system, we can use the pseudo-inverse or even better, the singular value decomposition to get the least squares solution. Our simple linear calibration approach has several advantages. It's easy to formulate, has a closed form solution, and often provides really good initial points for non-linear calibration approaches.

However, this simple approach has some disadvantages. One disadvantage of solving for \mathbf{P} , is that we do not directly get the intrinsic and extrinsic camera parameters. Furthermore, our linear model does not take into account complex phenomena, such as radial and tangential distortion. Finally, since we are solving via the linear least squares method, we cannot impose constraints on our solution, such as requiring the focal length to be non-negative.

The camera projection matrix \mathbf{P} by itself, is useful for projecting 3D points into 2D, but it has several drawbacks. It does not tell you the camera's pose and it does not tell you about the camera's internal geometry. Fortunately, we can factorize \mathbf{P} into intrinsic parameter matrix \mathbf{K} and extrinsic rotation parameters \mathbf{R} and translation parameters \mathbf{t} , using a linear algebra operation known as the RQ factorization. Let us see how we perform this factorization.

3.2 RQ Factorization

First, we alter the representation of \mathbf{P} to be a function of the camera center \mathbf{C} . \mathbf{C} is the point that projects to zero when multiplied by \mathbf{P} . We multiply \mathbf{K} into the matrix to form two sub-matrices, \mathbf{KR} and $-\mathbf{KRC}$. We will refer to the combination of \mathbf{K} and \mathbf{R} as the \mathbf{M} matrix. We can now express our projection matrix \mathbf{P} as

$$\mathbf{P} = [\mathbf{M} | -\mathbf{MC}] \quad (11)$$

From here, we use the fact that any square matrix can be factored into an upper triangular matrix \mathbf{R} and an orthogonal basis to decompose \mathbf{M} into upper triangular \mathbf{R} and orthogonal basis \mathbf{Q} . In linear algebra, this procedure is known as RQ factorization, which is a variant of the more commonly referred to QR factorization. In QR factorization, we have the orthogonal \mathbf{Q} first and then the upper triangular \mathbf{R} .

Remark 3.1. Note here that the \mathbf{R} and the output of RQ factorization, is a different variable than our rotation matrix \mathbf{R} .

Let us now see how we can use the output of RQ factorization of the matrix \mathbf{M} to retrieve \mathbf{K} , \mathbf{R} , and \mathbf{t} by aligning these two expressions. The intrinsic calibration matrix \mathbf{K} is the output \mathbf{R} of the RQ factorization of \mathbf{M} . The rotation matrix \mathbf{R} is the orthogonal basis \mathbf{Q} . Finally, we can extract the translation vector directly from \mathbf{K} in the last column of the \mathbf{P} matrix.

RQ factorization is a great tool to compute \mathbf{K} , \mathbf{R} , and \mathbf{t} from the camera \mathbf{P} matrix. However, some mathematical assumptions need to be performed to guarantee a unique solution for these matrices.

Monocular camera calibration is a well-established tool that has excellent implementations in C++, Python and MATLAB. OpenCV camera calibration see `CameraCalibration2`

3.3 Summary

So, to summarize. In this section, we discussed a method to identify the camera projection matrix \mathbf{P} . This method is known as camera calibration. We saw that the matrix \mathbf{P} can be factored into the following components

- Camera intrinsic matrix \mathbf{K}
- Camera extrinsic parameters \mathbf{R} and \mathbf{t} ,

This can be done through RQ factorization.

3.4 Questions

3.5 Assignments

4 Visual Depth Perception

Self-driving cars require accurate depth perception for the safe operation of our autonomous vehicles. If we do not know how far away the cars are in front of us, how can you avoid them while driving? Lidar and Radar sensors are

usually thought of as the primary 3D sensors available for perception tasks. However, we can get depth information from two or more cameras using multi-view geometry. Specifically, we will be describing the process of getting depth from two axis aligned cameras a setup known as the stereo cameras.

In this section, we will cover the geometry of the stereo sensor as well as how to derive the 3D coordinates of a point given its projection onto two images of the stereo sensor.

Remark 4.1. Some History

Stereopsis, the process of stereo vision, was first described by Charles Wheatstone back in 1838. He recognized that because each eye views the visual world from a slightly different horizontal position that each eye's image differs from the other. Objects at different distances from the eye project images into the two eyes that differ in their horizontal position giving depth cues of horizontal disparity that are also known as binocular disparity. However, historical evidence suggests that stereopsis was discovered much earlier than this. In fact some drawings by Leonardo da Vinci depict accurate geometry of depth through stereopsis. Up to the 19th century, the phenomenon of stereopsis was primarily used for entertainment. Anaglyphs were used to provide a stereoscopic 3D effect when viewed with 2D color glass, where each lens employs different chromatically opposite colors, usually red and cyan. Nowadays, we use stereopsis with complex algorithms to derive depth from two images using a similar concept to Da Vincis drawings.

Let us now delve into the geometry of a stereo sensor. A stereo sensor is usually created by two cameras with parallel optical axes. To simplify the problem even more, most manufacturers align the cameras in 3D space so that the two image planes are aligned with only an offset in the x -axis. Given a known rotation and translation between the two cameras and a known projection of a point \mathbf{O} in 3D to the two camera frames resulting in pixel locations OL and OR respectively, we can formulate the necessary equations to compute the 3D coordinates of the point \mathbf{O} . In order to make our computation easier, we will state some assumptions. First, we assume that the two cameras used to construct the stereo sensors are identical. Second, we will assume that while manufacturing the stereo sensor, we tried as hard as possible to keep the two cameras optical axes aligned. Let's now define some important parameters of the stereo sensor. The focal length f is the distance between the camera center and the image plane. Second, the baseline is defined as the distance along the shared x -axis between the left and right camera centers. By defining a baseline to represent the transformation between the two camera coordinate frames, we are assuming that the rotation matrix is identity and there is only a non-zero x component in the translation vector. The \mathbf{R} and \mathbf{T} transformation therefore boils down to a single baseline parameter b .

Before proceeding, we will project the previous figure to bird's eye view for easier visualization. Now, let's define the quantities we would like to compute. We want to compute the x and z coordinates of the point \mathbf{O} with respect to the left camera frame. The y coordinate can be estimated easily after the x and z coordinates are computed. Remember, we are given the baseline, focal length, and the coordinates of the projection of the point \mathbf{O} onto the left and right image planes. We can see two similar triangles formed by the left camera measurement as follows. The triangle formed by the depth z and the position x is similar to the triangle formed by the focal length f and the left measurement x_L . From this similarity we can construct the equation

$$\frac{Z}{f} = \frac{X}{x_L} \quad (12)$$

The same can be done for the right measurements but with the offset for the baseline included. In this case, the two triangles are defined by:

$$\frac{Z}{f} = \frac{X - b}{x_R} \quad (13)$$

Similarly, we can get a second equation relating Z to X via the right camera parameters in measurements.

From these two equations, we can now derive the 3D coordinates of the point \mathbf{O} . We define the disparity d to be the difference between the image coordinates of the same pixel in the left and right images. We can easily transform between image and pixel coordinates using the X and Y offsets u_0 and v_0 . We then use the two equations from the similar triangle relations to solve for the value of Z as follows. From there we use the value of z to compute X with the following expression.

$$X = \frac{Zx_L}{f} \quad (14)$$

Finally, we can repeat the process in the Y direction with the same derivation to arrive at the following expression for Y .

$$Y = \frac{Zy_L}{f} \quad (15)$$

The three components of the point position are now explicitly available from the two sets of pixel measurements available to us. Now that we have established the equations needed for 3D coordinate computation from the stereo sensor, two problems arise to be able to perform this computation. First, we need to compute the focal length baseline and x and y offsets. That is, we need to calibrate the stereo camera system. Second, we need to find the correspondence

between each left and right image pixel pair to be able to compute their disparity. Fortunately, the calibration problem can be solved using stereo camera calibration. This is an extension of the monocular process we discussed in section 3, for which well-established implementations are available. The correspondence problem however, requires specialized algorithms to efficiently perform the matching and compute the disparity between left and right image pixels, which we'll discuss further in the next video. The output depth from stereopsis suffers from some limitations particularly as points move further away from the stereo camera. However, given a good disparity estimation algorithm, the output is still useful for self-driving cars as a dense source of depth information and closer range which exceeds the density we can get from common Lidar sensors.

4.1 Summary

To summarize this section, we discussed the equations required to estimate 3D coordinates of a pixel given the geometric transformation between the two cameras sensors and the disparity between pixels. In the next section we will learn more about disparity generating algorithms and show full examples on how to compute that disparity from a stereo image pair using Python and OpenCV.

4.2 Questions

4.3 Assignments

5 Compute Disparity

So far we have learned the essential equations to extract 3D information from a stereo pair. However, we were faced with some unknown parameters that we have to estimate.

In this section we will learn how to estimate these missing parameters such as the disparity through stereo matching. We will also learn that efficient disparity estimation is possible due to epipolar constraints.

Remark 5.1. Epipolar Geometry

See the wikipedia article at https://en.wikipedia.org/wiki/Epipolar_geometry for epipolar geometry. See also the OpenCV article at https://docs.opencv.org/3.4.3/da/de9/tutorial_py_epipolar_geometry.html and the following videos on Coursera

- <https://www.coursera.org/lecture/robotics-perception/epipolar-geometry-i-oNH09>

- <https://www.coursera.org/lecture/robotics-perception/epipolar-geometry-ii-WRyoL>
- <https://www.coursera.org/lecture/robotics-perception/epipolar-geometry-iii-Bwk0d>

Recall from section 4 that we identified two primary issues with the visual depth estimation from stereo images.

- The camera parameters focal length f baseline b
- The camera pixel centers u_0, v_0

These need to be estimated from stereo camera calibration. Similar to monocular camera calibration, stereo calibration is a well-studied problem with lots of user-friendly free software capable of performing in. In this section we will be targeting the second problem, mainly stereo matching to compute disparities. As a reminder, disparity is the difference in the image location of the same 3D point as observed by two different cameras.

5.1 Disparity Computation Algorithm

Before we begin let us explain what do we mean with the term disparity

Definition 5.1. Disparity

With the term disparity we mean the difference in image location of the same 3D point under perspective to two different cameras.

To compute the disparity we need to be able to find the same point in the left and right stereo camera images. In other words, we need to find \mathbf{x}_R for each corresponding \mathbf{x}_L . This problem is known as the **stereo correspondence problem**.

Remark 5.2. The Stereo Correspondence Problem

The simplest solution, however naive one, for this problem is an exhaustive search i.e we search the whole right image for every pixel in the left image. Such a solution is extremely inefficient and will usually not perform in real time to be used on self-driving cars. It is also unlikely to succeed as many pixels will have similar local characteristics, making it difficult to match them correctly. Luckily for us, we can use stereo geometry to constrain our search problem from 2D over the entire image space to a 1D line.

Let us revisit the stereo camera setup and see why such a simplification is valid. We have already determined, how a single point is projected to both

cameras. Now, let's move our 3D point along the line connecting it with the left camera's center. Its projection on the left camera image plane does not change. However, what can you notice about the projection on the right camera plane, the projection moves along the horizontal line. This is called an **epipolar line**, see Figure 1.




Fig. 1: Epipolar line.

The epipolar line follows directly from the fixed lateral offset and image plane alignment of the two cameras in a stereo pair. We can constrain our correspondence search to be along the epipolar line, reducing the search from 2D to 1D. One thing to note is that horizontal epipolar lines only occur if the optical axes of the two cameras are parallel. In the case of non parallel optical axis, the epipolar lines are skewed. In such cases we will have to resort to multiple view geometry rather than the stereo equations we have developed.

In the case of two calibrated cameras, such as our stereo camera, a skewed epipolar line is not a huge problem. In fact, we can work the optical axis to be parallel through a process known as **stereo rectification**. After rectification we arrive back to our horizontal epipolar line. We will not go through how to perform rectification as implementations are available in standard computer vision packages such as OpenCV and MATLAB.

Remark 5.3. Stereo Rectification

See the following article on wikipedia about image rectification https://en.wikipedia.org/wiki/Image_rectification. Also OpenCV base stereo image calibration <https://sourishghosh.com/>

`2016/stereo-calibration-cpp-opencv/.`

Let us go over our first basic stereo algorithm.

1. For each epipolar line take a pixel on this line in the left image
2. Compare these left image pixels to every pixel in the right image on the same epipolar line.
3. Pick the pixel that has minimum cost. For example, a very simple cost here can be the squared difference in pixel intensities.
4. Compute disparity d by subtracting the right image location from the left one

Stereo matching is a very well-studied problem in computer vision. Many more complex costs and search regions can be defined, which attempts to improve either computational efficiency or disparity accuracy. There are a wide range of approaches including both local and global methods, which differ in the main image region considered when identifying correspondences and computing disparities. As with most problems in computer vision, the stereo vision algorithms are evaluated on a public benchmark. The most famous of which is the Middlebury stereo benchmark, <http://vision.middlebury.edu/stereo/eval3/>. If you are interested, many of the top-performing stereo matching algorithms have results published there and have code available too.

5.2 Summary

5.3 Questions

5.4 Assignments

References

- [1] https://en.wikipedia.org/wiki/Pinhole_camera
- [2] https://boofcv.org/index.php?title=Tutorial_Camera_Calibration
- [3] https://docs.opencv.org/3.4.3/dc/dbb/tutorial_py_calibration.html